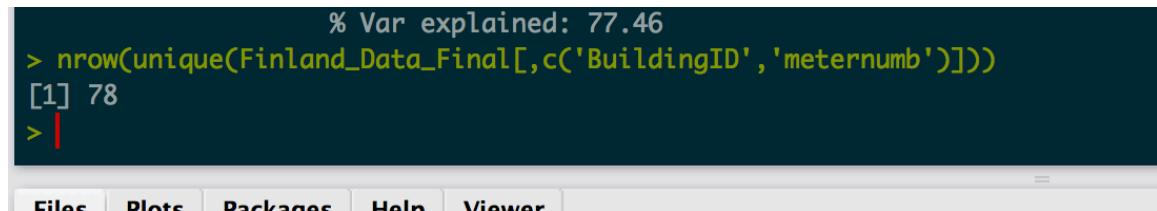


Part 2:

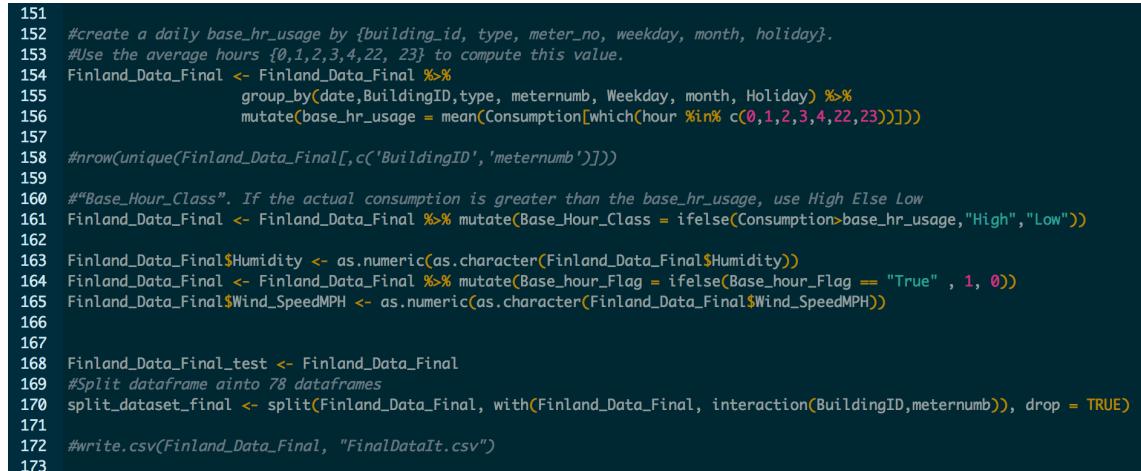
- You will need to create a daily base_hr_usage by {building_id, type, meter_no, weekday, month, holiday}.
- Use the average hours {0,1,2,3,4,22, 23} to compute this value.
- Create a column base_hr_usage and input the calculated value in that column. Note that you will have one value for each hour for the whole day.
- Create another column called “Base_Hour_Class”. If the actual consumption is greater than the base_hr_usage, use High Else Low

→ We can find number of unique combinations of Building id and meter number:

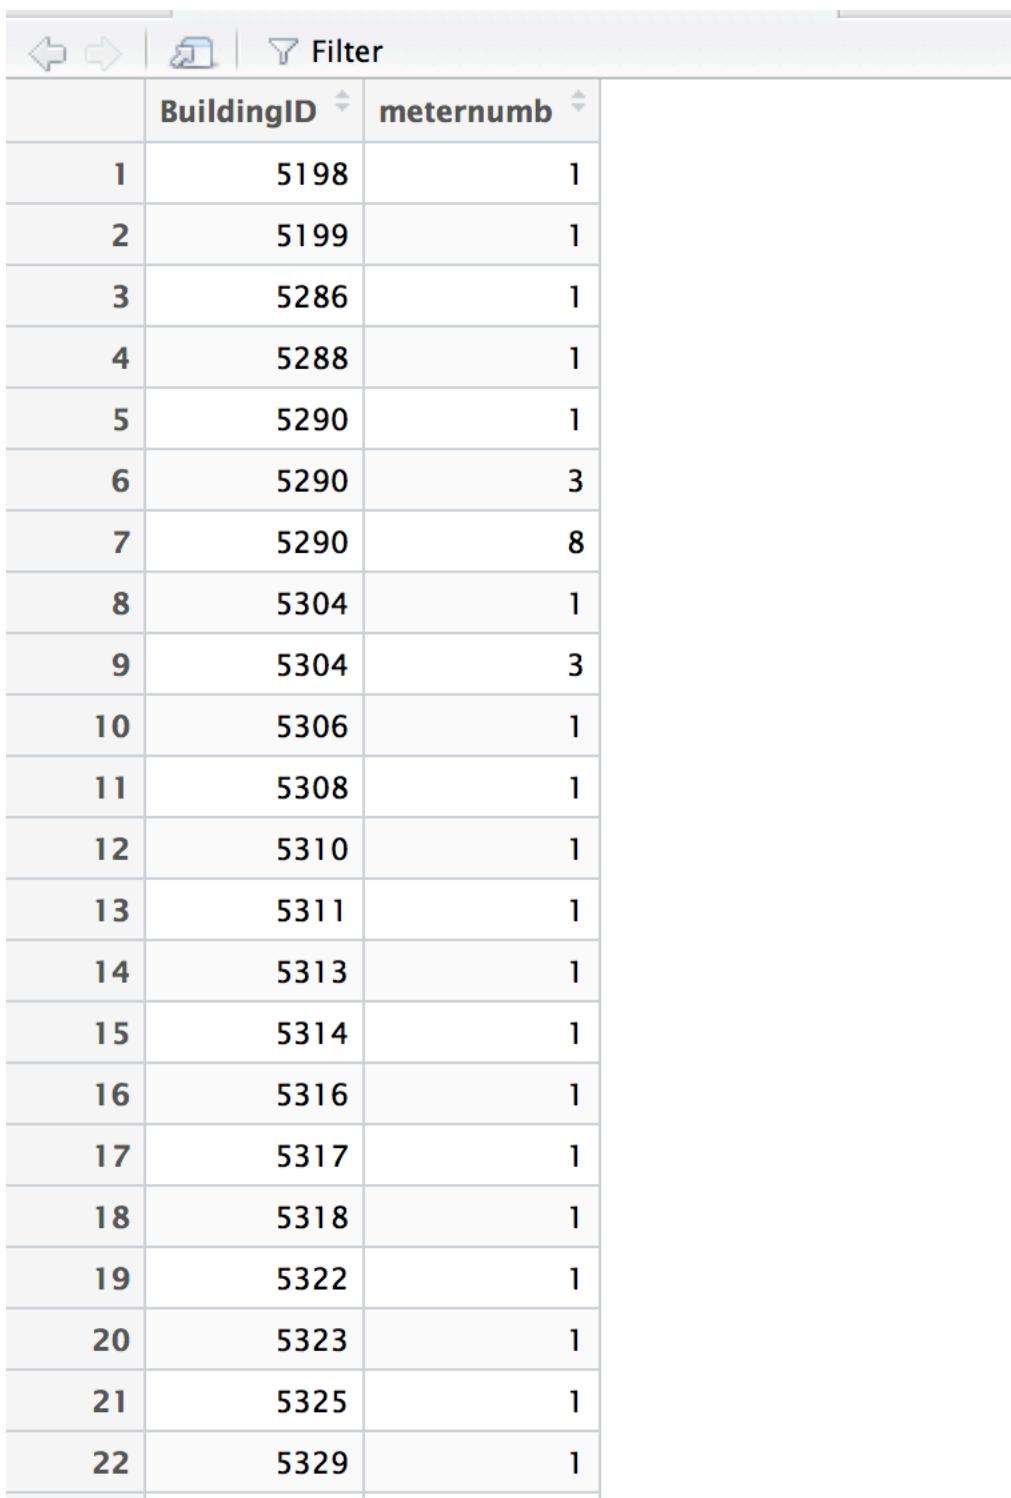
```
nrow(unique(Finland_Data_Final[,c('BuildingID','meternumb')]))
```



```
% Var explained: 77.46
> nrow(unique(Finland_Data_Final[,c('BuildingID','meternumb')]))
[1] 78
> |
```



```
151
152 #create a daily base_hr_usage by {building_id, type, meter_no, weekday, month, holiday}.
153 #Use the average hours {0,1,2,3,4,22, 23} to compute this value.
154 Finland_Data_Final <- Finland_Data_Final %>%
155   group_by(date,BuildingID,type, meternumb, Weekday, month, Holiday) %>%
156   mutate(base_hr_usage = mean(Consumption[which(hour %in% c(0,1,2,3,4,22,23))]))
157
158 #nrow(unique(Finland_Data_Final[,c('BuildingID', 'meternumb')]))
159
160 ##"Base_Hour_Class". If the actual consumption is greater than the base_hr_usage, use High Else Low
161 Finland_Data_Final <- Finland_Data_Final %>% mutate(Base_Hour_Class = ifelse(Consumption>base_hr_usage,"High","Low"))
162
163 Finland_Data_Final$Humidity <- as.numeric(as.character(Finland_Data_Final$Humidity))
164 Finland_Data_Final <- Finland_Data_Final %>% mutate(Base_hour_Flag = ifelse(Base_hour_Flag == "True" , 1, 0))
165 Finland_Data_Final$Wind_SpeedMPH <- as.numeric(as.character(Finland_Data_Final$Wind_SpeedMPH))
166
167
168 Finland_Data_Final_test <- Finland_Data_Final
169 #Split dataframe ainto 78 dataframes
170 split_dataset_final <- split(Finland_Data_Final, with(Finland_Data_Final, interaction(BuildingID,meternumb)), drop = TRUE)
171
172 #write.csv(Finland_Data_Final, "FinalDataIt.csv")
173
```



A screenshot of a Microsoft Access query results grid. The grid has two columns: "BuildingID" and "meternumb". The "BuildingID" column contains values from 1 to 22, and the "meternumb" column contains values 1, 1, 1, 1, 1, 3, 8, 1, 3, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1. The grid includes standard Access navigation buttons (Back, Forward, New, Save) and a Filter button at the top.

	BuildingID	meternumb
1	5198	1
2	5199	1
3	5286	1
4	5288	1
5	5290	1
6	5290	3
7	5290	8
8	5304	1
9	5304	3
10	5306	1
11	5308	1
12	5310	1
13	5311	1
14	5313	1
15	5314	1
16	5316	1
17	5317	1
18	5318	1
19	5322	1
20	5323	1
21	5325	1
22	5329	1

23	5332	1
24	5333	1
25	5335	1
26	5340	1
27	5351	1
28	5352	1
29	5355	1
30	5356	1
31	5357	1
32	5358	1
33	28096	1
34	28096	4
35	28108	3
36	28161	1
37	28162	1
38	28168	1
39	28168	4
40	30602	1
41	75689	1
42	75690	1
43	75691	1
44	75692	1

	BuildingID	meternumb
45	75693	1
46	75694	1
47	75696	1
48	75696	2
49	75697	1
50	75698	1
51	75699	1
52	75700	1
53	75701	1
54	75702	1
55	75703	1
56	75704	1
57	75705	1
58	75706	1
59	75707	1
60	75708	1
61	75709	1
62	75710	1
63	75711	1
64	75712	1
65	75713	1
66	75714	1
--	-----	-

55	75705	1
57	75705	1
58	75706	1
59	75707	1
60	75708	1
61	75709	1
62	75710	1
63	75711	1
64	75712	1
65	75713	1
66	75714	1
67	75715	1
68	75716	1
69	75717	1
70	75718	1
71	75719	1
72	75720	1
73	75721	1
74	77662	2
75	81909	1
76	82254	1
77	83427	1
78	84681	1

- Prediction

Neural Network Prediction:

Using the `glm()` function instead of the `lm()` this will become useful later when cross validating the linear model.

```
2 set.seed(123)
3 data2 <- split_dataset_final[[1]]
4 data2 <- data2[,c(18,6,8,9,10,12,19,21,22,16,26)]
5 data2 <- as.data.frame(data2)
6 data_original <- data2
7 apply(data2,2,function(x) sum(is.na(x)))
8 normalize<-function(x)(return((x-min(x))/(max(x)-min(x))))
9 data1<-as.data.frame(lapply(data2[,c(1:9,11)],normalize))
10 data3 <- as.data.frame(data2$Normalized_Consumption)
11 colnames(data3) <- c("Normalized_Consumption")
12 data <- cbind(data1, data3)
13
14 index <- sample(1:nrow(data),round(0.75*nrow(data)))
15 index_original <- index
16 train <- data[index,]
17 test <- data[-index,]
18 lm.fit <- glm(Normalized_Consumption~TemperatureF+hour+dayOfWeek+Weekday+month
19                 +Holiday+Dew_PointF+Sea_Level_PressureIn
20                 +VisibilityMPH+WindDirDegrees
21                 , data=train)
22 summary(lm.fit)
23 pr.lm <- predict(lm.fit,test)
24 MSE.lm <- sum((pr.lm - test$Normalized_Consumption)^2)/nrow(test)
25
26
```

```
> MSE.lm
[1] 0.002173556524
```

Preparing to fit the neural network

It is good practice to normalize your data before training a neural network. Usually scaling in the intervals [0,1] or [-1,1] tends to give better results. We therefore scale and split the data before moving on:

`scale` returns a matrix that needs to be coerced into a `data.frame`.

```

#Preparing to fit the neural network
maxs <- apply(data, 2, max)
mins <- apply(data, 2, min)

scaled <- as.data.frame(scale(data, center = mins, scale = maxs - mins))
train_ <- scaled[index,]
test_ <- scaled[-index,]

```

Usually, if at all necessary, one hidden layer is enough for a vast numbers of applications. As far as the number of neurons is concerned, it should be between the input layer size and the output layer size, usually 2/3 of the input size.

2 hidden layers will be used

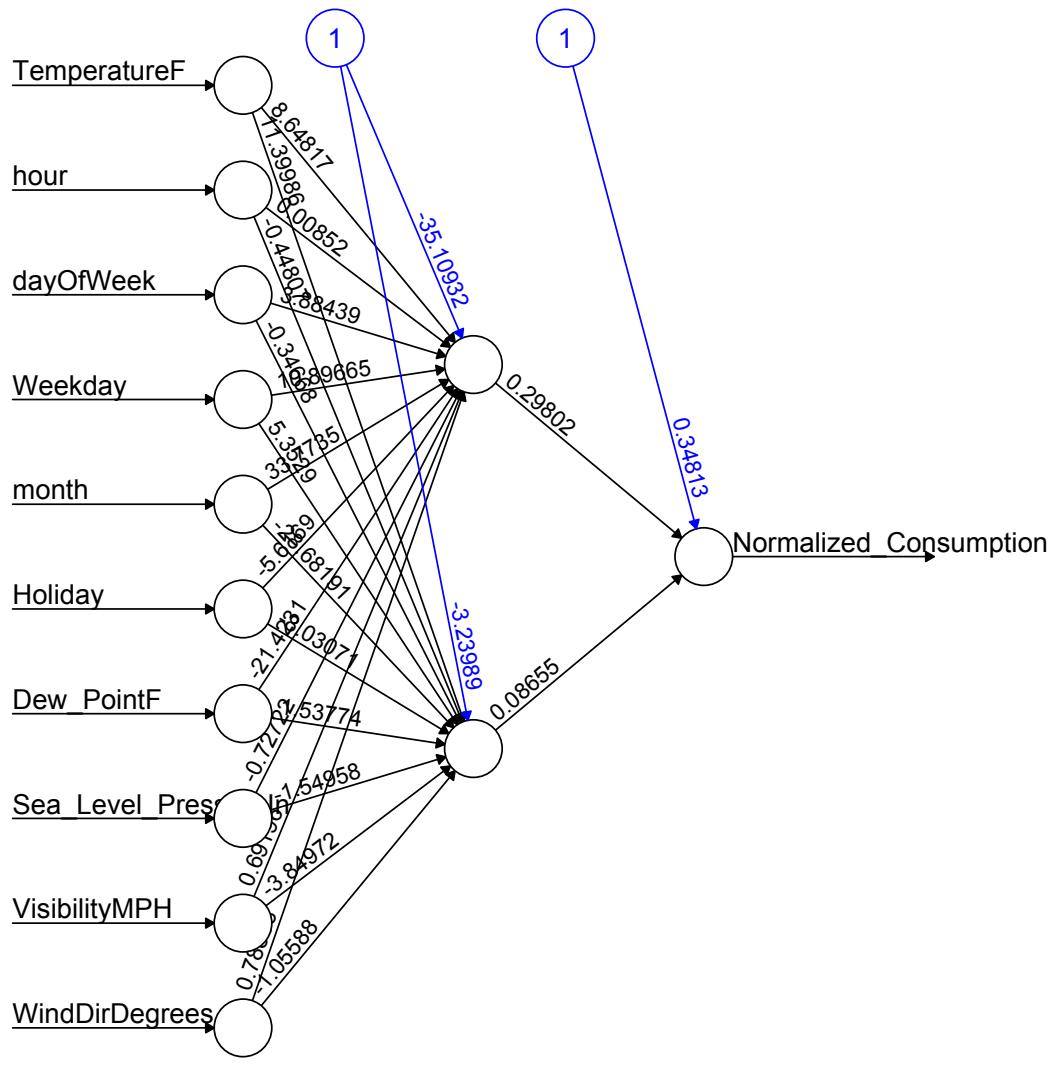
The input layer has 10 inputs, the two hidden layers have 2 neurons and the output layer has, of course, a single output since we are doing regression.

fit the net:

```

1 n <- names(train_)
2 f <- as.formula(paste("Normalized_Consumption ~", paste(n[!n %in% "Normalized_Consumption"], collapse = " + ")))
3 nn <- neuralnet(f,data=train_,hidden=2,linear.output=T)
4 nn <- neuralnet(f,data=train_,hidden=2,linear.output=T)
5 plot(nn)
6 print(nn)

```



Error: 15.512083 Steps: 15082

The black lines show the connections between each layer and the weights on each connection while the blue lines show the bias term added in each step.

Predicting Normalized Consumption using the neural network

```
#Predicting Normalized Consumption using the neural network

pr.nn <- compute(nn,test_[,1:10])
pr.nn_ <- pr.nn$net.result*(max(data$Normalized_Consumption)-
                           min(data$Normalized_Consumption))+min(data$Normalized_Consumption)
test.r <- (test_$Normalized_Consumption)*(max(data$Normalized_Consumption)-
                                             min(data$Normalized_Consumption))+min(data$Normalized_Consumption)
MSE.nn <- sum((test.r - pr.nn_)^2)/nrow(test_)
#we then compare the two MSEs

print(paste(MSE.lm,MSE.nn))
```

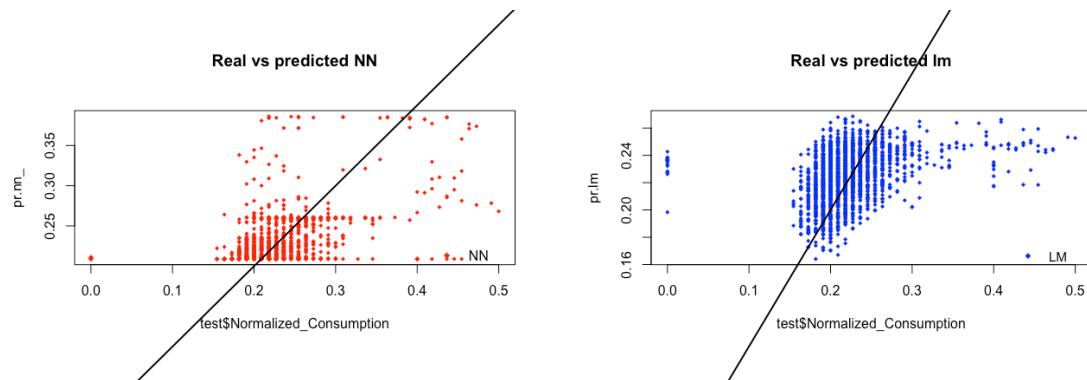
```
> print(paste(MSE.lm,MSE.nn))
[1] "0.00217355652434915 0.00185239396341829"
```

Apparently the net is doing a better work than the linear model at prediction.

Below, after the visual plot, we are going to perform a fast cross validation in order to be more confident about the results.

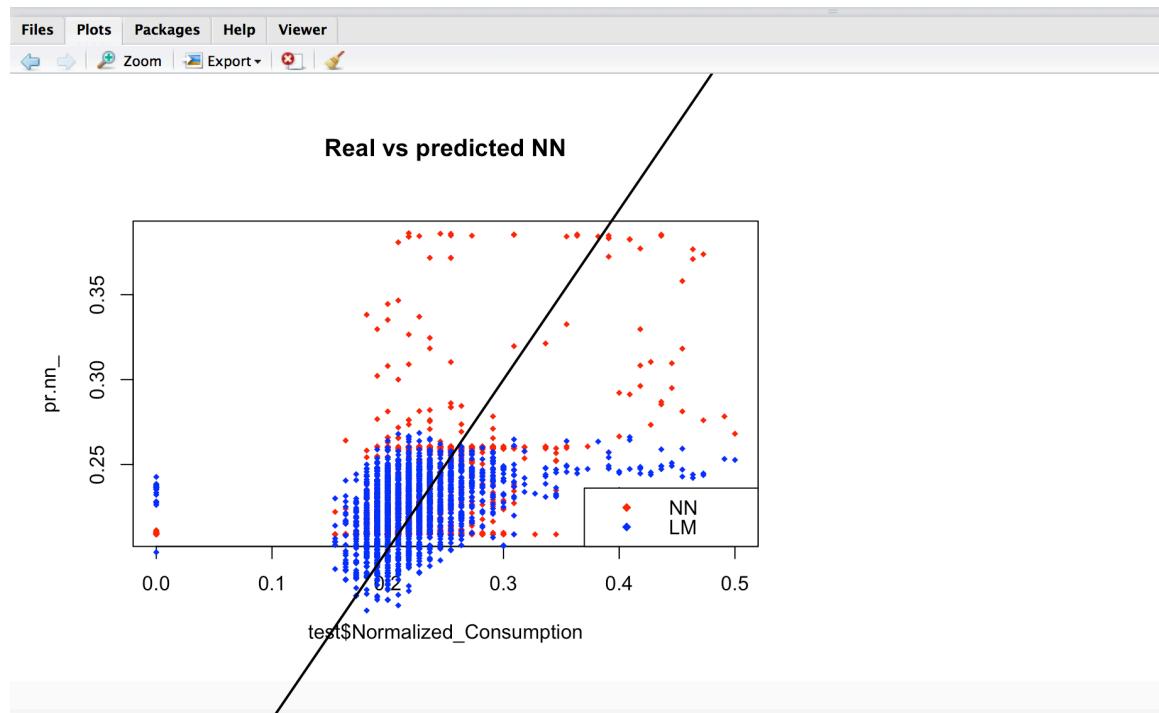
```
par(mfrow=c(1,2))
plot(test$Normalized_Consumption,pr.nn_,col='red',main='Real vs predicted NN',pch=18,cex=0.7)
abline(0,1,lwd=2)
legend('bottomright',legend='NN',pch=18,col='red', bty='n')

plot(test$Normalized_Consumption,pr.lm,col='blue',main='Real vs predicted lm',pch=18, cex=0.7)
abline(0,1,lwd=2)
legend('bottomright',legend='LM',pch=18,col='blue', bty='n', cex=.95)
```



By visually inspecting the plot we can see that the predictions made by the neural network are (in general) more concentrated around the line (a perfect alignment with the line would indicate a MSE of 0 and thus an ideal perfect prediction) than those made by the linear model.

A perhaps more useful visual comparison is plotted below:



Prediction:

```
#Prediction

pr.nn_ <- as.data.frame(pr.nn_)
colnames(pr.nn_) <- c("Predicted value")
data2_test <- data_original[-index_original,]
Final_prediction_nn <- cbind(data2_test, pr.nn_)

#calculating mean square value
rmse(pr.nn_,data2_test$Normalized_Consumption )
rmse <- c("RMS", rmse(pr.nn_,data2_test$Normalized_Consumption))
mae <- c("MAE", mae(pr.nn_,data2_test$Normalized_Consumption))

# Function that returns Mean Absolute Percentage Error

mape <- function(error) {
  mean(abs(error/Neuraldataset$kWh) * 100)
}

mape <- c("MAPE", mape(pr.nn_,data2_test$Normalized_Consumption))
```

```
standard_deviation <- sd(Final_prediction_nn[,10])
```

```
> standard_deviation
```

```
[1] 0.04907350636
```

```
> Final_prediction_nn <- Final_prediction_nn %>% mutate(Outlier_Tag = ifelse(residual >= (2*standard_deviation) , 1, 0))
```

```
> unique(Final_prediction_nn$Outlier_Tag)
```

```
[1] 0 1
```

	TemperatureF	hour	dayOfWeek	Weekday	month	Holiday
2	33.8	1		2	1	1
3	33.8	2		2	1	1
5	35.6	4		2	1	1
10	37.4	9		2	1	1
16	37.4	15		2	1	1
17	37.4	16		2	1	1
	Dew_PointF	Sea_Level_PressureIn	VisibilityMPH			
2	32.0		29.42		1.9	
3	33.8		29.39		3.7	
5	33.8		29.33		5.0	
10	35.6		29.30		4.3	
16	35.6		29.33		1.9	
17	37.4		29.33		2.5	
	Normalized_Consumption	WindDirDegrees	Predicted	value		
2	0.2272727273		150	0.2290355289		
3	0.2363636364		150	0.2267607404		
5	0.2272727273		150	0.2272411944		
10	0.2363636364		180	0.2290167902		
16	0.2272727273		190	0.2306024999		
17	0.2181818182		190	0.2297830560		
	residual	Outlier_Tag				
2	-0.00176280167322	0				
3	0.00960289597545	0				
5	0.00003153285363	0				
10	0.00734684614090	0				
16	-0.00332977258711	0				
17	-0.01160123784468	0				

#Calculating mape, mae, rmse

```
#calculating mean square value
rmse<- rmse(pr.nn_,data2_test$Normalized_Consumption )
rmse <- c("RMS", rmse(pr.nn_,data2_test$Normalized_Consumption))
mae <- c("MAE", mae(pr.nn_,data2_test$Normalized_Consumption))

# Function that returns Mean Absolute Percentage Error

mape <- abs((Final_prediction_nn$residual/Final_prediction_nn$Normalized_Consumption)*100)
mape <- mean(mape)
mape <- c("MAPE", mape)
neural_performance <- rbind(rmse, mae, mape, deparse.level = 0)
```

```
[1,] "RMS" "0.0430394465974911"
[2,] "MAE" "0.0270693978301965"
[3,] "MAPE" "Inf"
```

Regression Prediction

Fit a regression tree to the data set. create a training set, and fit the tree to the training data.

```
#fit a regression tree to the data set. create a training set, and fit the tree to the training data.
set.seed(123)
data_reg_tree <- split_dataset_final[[1]]
data_reg_tree <- data_reg_tree[,c(18,6,9,10,12,14,19,20,26,16)]

train = sample (1:nrow(data_reg_tree), nrow(data_reg_tree)/2)
tree.data = tree(Normalized_Consumption~.,data_reg_tree,subset=train)
summary (tree.data)
```

Regression tree:

```
tree(formula = Normalized_Consumption ~ ., data = data_reg_tree,
  subset = train)
```

Variables actually used in tree construction:

```
[1] "month"
[2] "Weekday"
[3] "hour"
[4] "WindDirDegrees"
[5] "TemperatureF"
[6] "Dew_PointF"
[7] "Humidity"
```

Number of terminal nodes: 15

Residual mean deviance: 0.001415731 = 5.790341 / 4090

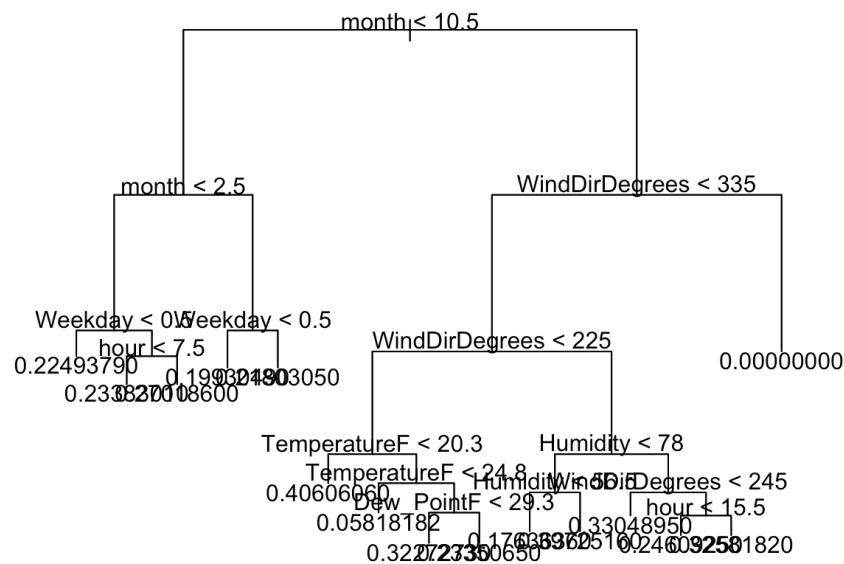
Distribution of residuals:

Min.	1st Qu.
-0.322727300	-0.018030520
Median	Mean
0.000151299	0.000000000
3rd Qu.	Max.
0.018333120	0.268181800

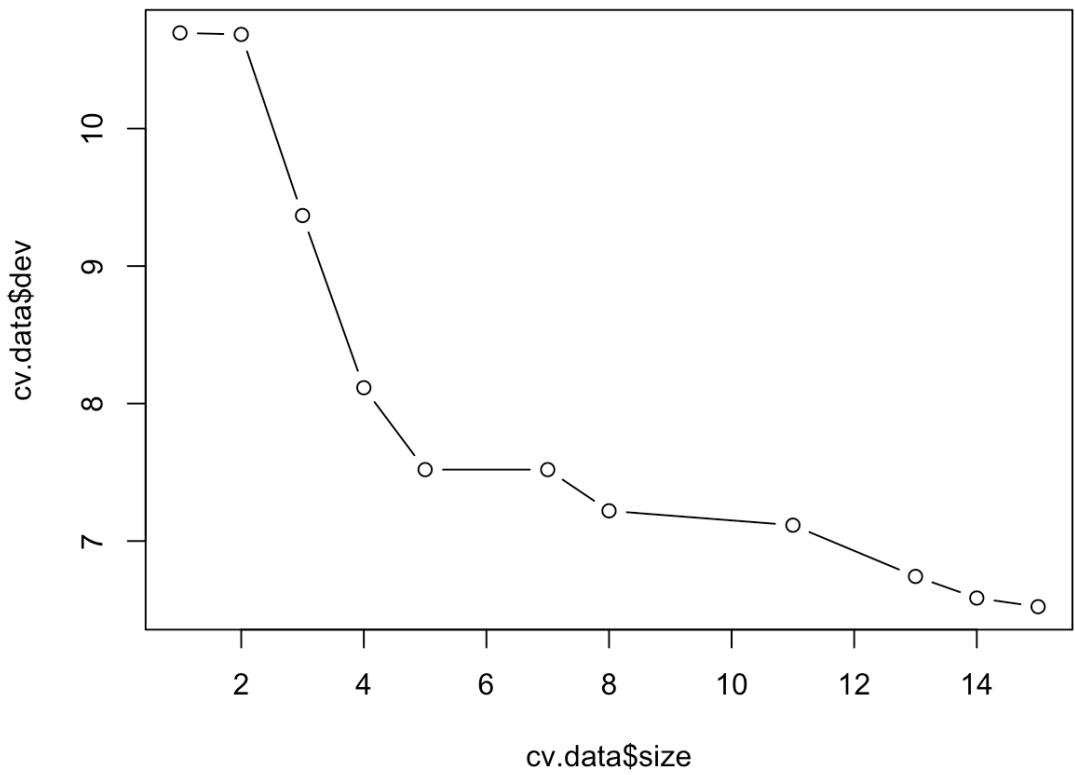
Only, 7 variables have been used to construct a tree

```
#Plotting a tree
```

```
#plot the tree.
plot(tree.data)
text(tree.data, pretty = 0)
#Cross-validation for Choosing Tree Complexity
cv.data = cv.tree(tree.data)
plot(cv.data$size, cv.data$dev, type='b')
#prune the tree, we could use the prune.tree() function:
prune.data =prune.tree(tree.data, best = 5)
plot(prune.data)
text(prune.data, pretty = 0)
#In keeping with the cross-validation results, we use the unpruned tree to make predictions on the test set.
yhat=predict (tree.data, newdata =data_reg_tree [-train,])
data.test=data_reg_tree [-train,"Normalized_Consumption"]
data.test1 <- as.double(data.test[['Normalized_Consumption']])
plot(yhat,data.test1)
abline (0,1)
```

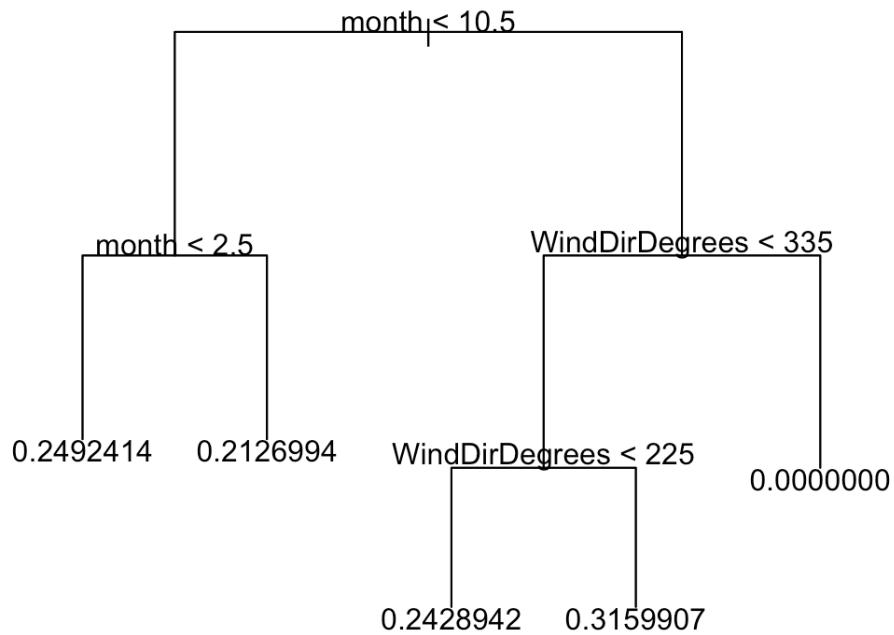


```
#Cross-validation for Choosing Tree Complexity
```

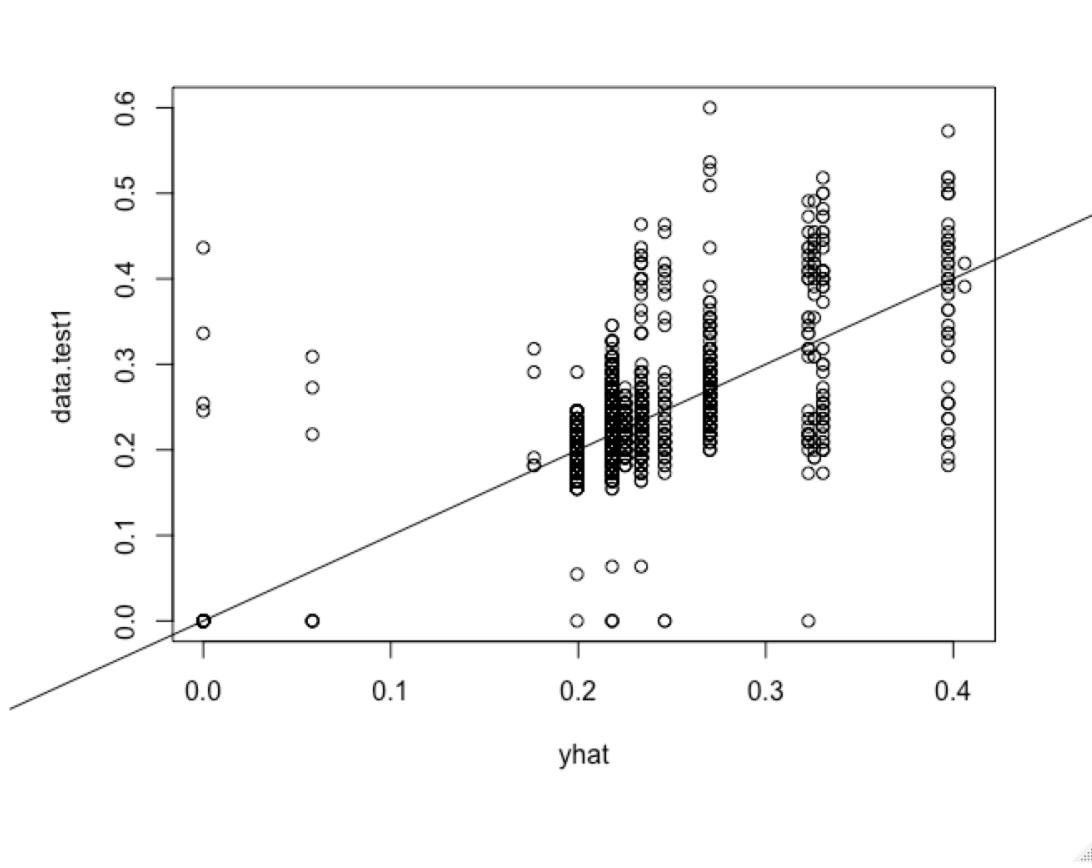


See whether the pruning will improve the performance:

prune the tree, we could use the `prune.tree()` function:



In keeping with the cross-validation results, we use the unpruned tree to make predictions on the test set.

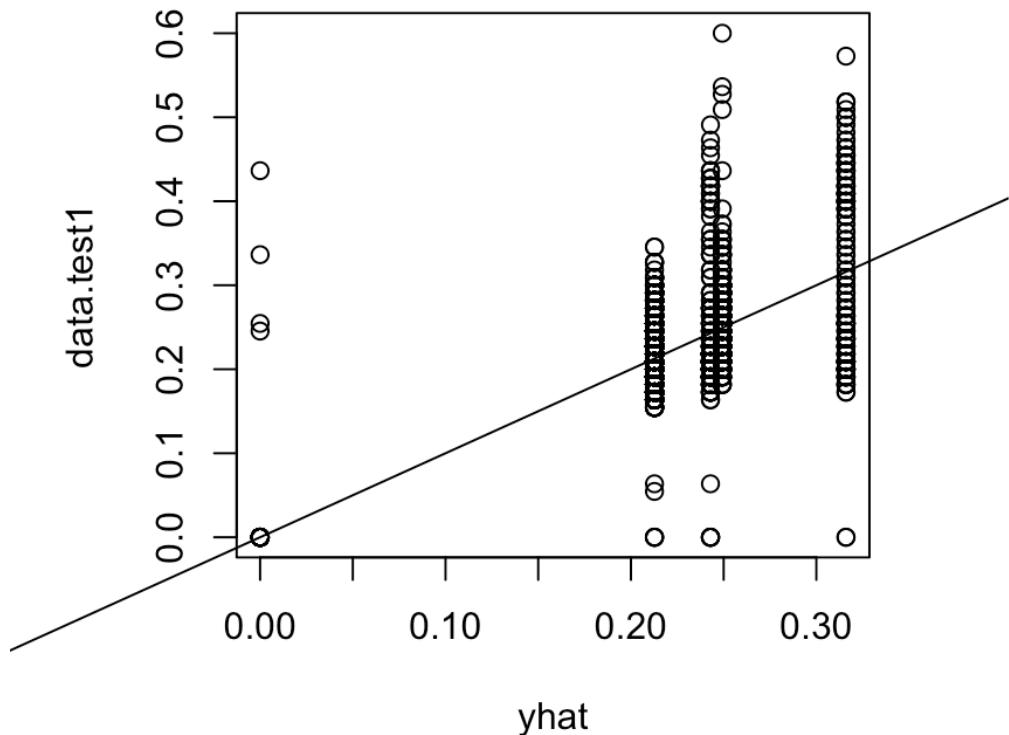


```
#MSE
mean((yhat -data.test)^2)
accuracy_1 <- accuracy(yhat, data.test1)
```

```
> mean((yhat -data.test)^2)
[1] 0.001599257638
> accuracy_1 <- accuracy(yhat, data.test1)
> accuracy_1
      ME      RMSE      MAE      MPE      MAPE
Test set 0.0001223930285 0.03999071939 0.02517541932 -Inf Inf
```

After removing the least important splits, the MSE only increases a little

```
#After removing the least important splits, the MSE only increases a little
yhat=predict (prune.data, newdata =data_reg_tree [-train,])
data.test=data_reg_tree [-train,"Normalized_Consumption"]
data.test1 <- as.double(data.test[['Normalized_Consumption']])
plot(yhat,data.test1)
abline (0,1)
#MSE
mean((yhat -data.test)^2)
accuracy_2 <- accuracy(yhat, data.test1)
```



```

> mean((yhat -data.test)^2)
[1] 0.001894604213
> accuracy_2 <- accuracy(yhat, data.test1)
> accuracy_2
      ME      RMSE      MAE      MPE      MAPE
Test set -0.00008241686817 0.04352705151 0.02824680772 -Inf Inf

```

So, mean is improved now

#Prediction

```

test <- as.data.frame( data_reg_tree [-train,variable_unique])
yhat_1 <- as.data.frame(yhat)
colnames(yhat_1) <- c("Predicted Value")
data.test_1 <- as.data.frame(data.test)
Final_prediction <- cbind(test,yhat_1,data.test_1)

```

```

> Final_prediction$residual <- Final_prediction$Normalized_Consumption -
+   Final_prediction_mn$`Predicted value`
> standard_deviation <- sd(Final_prediction$residual)
> standard_deviation
[1] 0.05694269232
> Final_prediction <- Final_prediction %>% mutate(Outlier_Tag = ifelse(residual >= (2*standard_deviation) , 1, 0))
> head(Final_prediction)
  month Weekday hour WindDirDegrees TemperatureF Dew_PointF Humidity Predicted Value
1     1       1     1           150        33.8      32.0     93  0.2492413617
2     1       1     2           150        33.8      33.8    100  0.2492413617
3     1       1     4           150        35.6      33.8     93  0.2492413617
4     1       1     6           180        37.4      35.6     93  0.2492413617
5     1       1     7           180        37.4      35.6     93  0.2492413617
6     1       1     9           180        37.4      35.6     93  0.2492413617
  Normalized_Consumption      residual Outlier_Tag
1          0.2272727273 -0.00176280167322      0
2          0.2363636364  0.00960289597545      0
3          0.2272727273  0.00003153285363      0
4          0.2272727273 -0.00174406295001      0
5          0.2454545455  0.01485204559471      0
6          0.2363636364  0.00658058033714      0
> |

```

Random Forest Prediction

Random forests improve predictive accuracy by generating a large number of bootstrapped trees (based on random samples of variables), classifying a case using each tree in this new "forest", and deciding a final predicted outcome by combining the results across all of the trees (an average in regression, a majority vote in classification).

```
1 # Random Forests
2 # Random forests improve predictive accuracy by generating a large number of bootstrapped trees
3
4 library(randomForest)
5 set.seed(123)
6 data_reg_random_forest <- split_dataset_final[[1]]
7 index_1 <- sample(1:nrow(data_reg_random_forest), round(0.75*nrow(data_reg_random_forest)))
8 index_original_1 <- index_1
9 train <- data_reg_random_forest[index_1,]
10 test <- data_reg_random_forest[-index_1,]
11
12 fit <- randomForest(Normalized_Consumption ~ TemperatureF+hour+Weekday+month
13                         +Dew_PointF +WindDirDegrees +Humidity, data = train)
14 print(fit) # view results
15 importance(fit) # importance of each predictor
16 prediction_random_forest <- predict(fit, test)
17 accuracy(prediction_random_forest, test$Normalized_Consumption)
18 #-----
```

Call:

```
randomForest(formula = Normalized_Consumption ~ TemperatureF + hour
+ Weekday + month + Dew_PointF + WindDirDegrees + Humidity, data =
train)
```

Type of random forest: regression

Number of trees: 500

No. of variables tried at each split: 2

Mean of squared residuals: 0.0008449698134

% Var explained: 68.7

```
> importance(fit) # importance of each predictor
  IncNodePurity
TemperatureF  1.950842558
hour         2.165137500
Weekday      1.160022237
month        3.363550462
Dew_PointF  1.776769290
WindDirDegrees 2.648233099
Humidity     1.608882317
```

```
data2_test <- test[,c(18,6,9,10,19,26,20,16)]
predicion_random_forest <- as.data.frame(predicion_random_forest)
Final_prediction_rr <- cbind(data2_test, predicion_random_forest)
Final_prediction_rr$residual <- Final_prediction_rr$Normalized_Consumption -
  Final_prediction_rr$predicion_random_forest
standard_deviation <- sd(Final_prediction_rr$residual)
Final_prediction_rr <- Final_prediction_rr %>%
  mutate(Outlier_Tag = ifelse(residual >= (2*standard_deviation) , 1, 0))

#calculating mean square value
accuracy(predicion_random_forest, test$Normalized_Consumption)
#-----
```

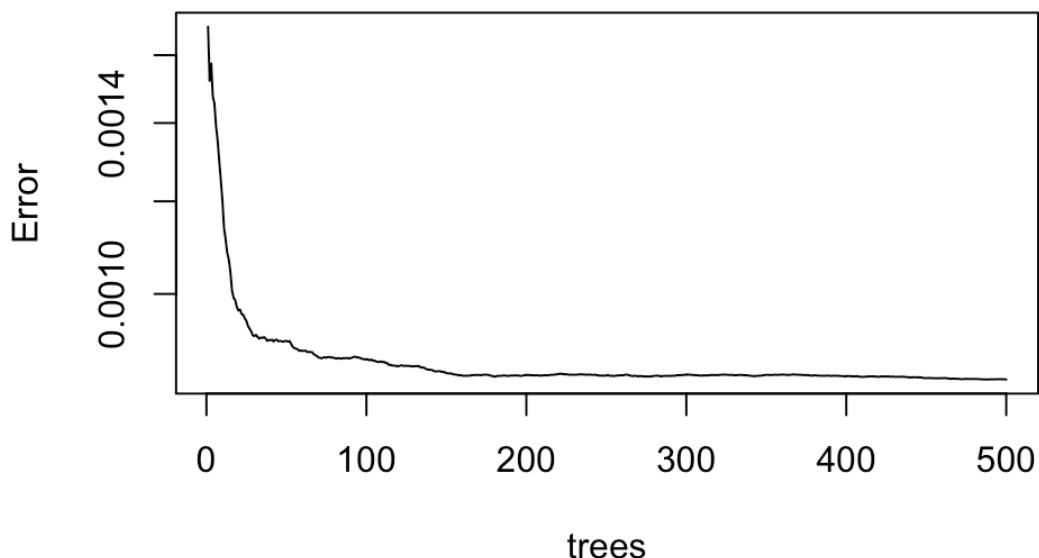
➤ **accuracy(predicion_random_forest,
Final_prediction_rr\$Normalized_Consumption)**

```
> head(Final_prediction_rr)
   TemperatureF hour Weekday month Dew_PointF WindDirDegrees Humidity
1          33.8    1      1     1       32.0           150       93
2          33.8    2      1     1       33.8           150      100
3          35.6    4      1     1       33.8           150       93
4          37.4    9      1     1       35.6           180       93
5          37.4   15      1     1       35.6           190       93
6          37.4   16      1     1       37.4           190      100
Normalized_Consumption predicion_random_forest      residual Outlier_Tag
1              0.2272727273      0.2318583389 -0.004585611631       0
2              0.2363636364      0.2290686577  0.007294978621       0
3              0.2272727273      0.2371340468 -0.009861319527       0
4              0.2363636364      0.2273727179  0.008990918438       0
5              0.2272727273      0.2317184482 -0.004445720887       0
6              0.2181818182      0.2338186263 -0.015636808137       0
> |
```

```
> randomforest_performance  
      ME        RMSE       MAE       MPE MAPE  
Test set -0.001211284622 0.04832920999 0.01674909995 -Inf Inf
```

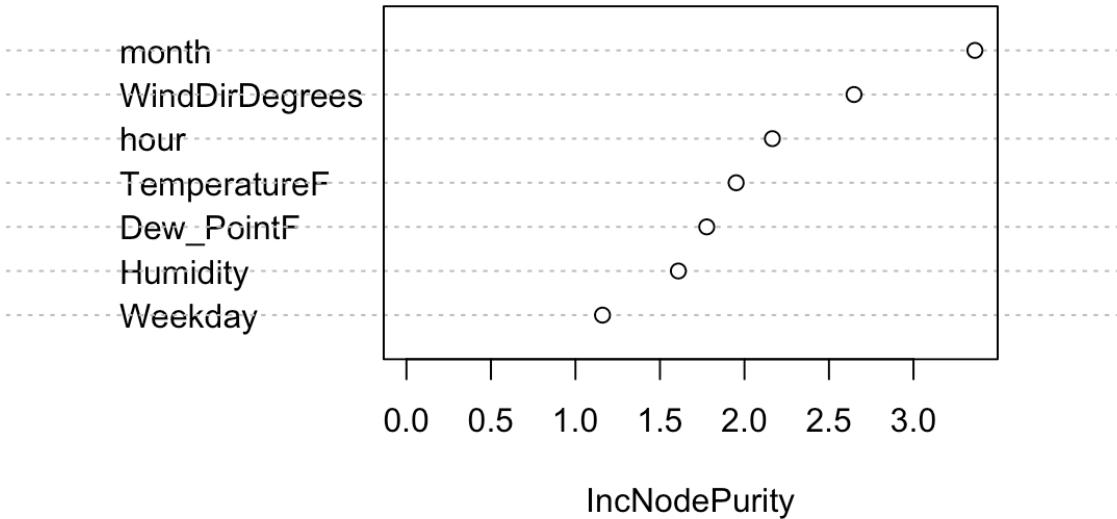
```
plot(fit, log="y")
```

fit



```
varImpPlot(fit)
```

fit



knn-Regression Prediction

```
1 install.packages("FNN")
2 library(FNN)
3 data2 <- split_dataset_final[[1]]
4 data2 <- data2[,c(18,6,8,9,10,12,19,21,22,16,26)]
5 data2 <- as.data.frame(data2)
6 data_original <- data2
7
8 train <- sample(1:nrow(data2),round(0.75*nrow(data2)))
9 traindata <- data2[train,]
10 testdata <- data2[-train,]
11
12 tree_knn<-knn.reg(traindata,y=traindata$Normalized_Consumption,testdata,k=3)
13 testdata$Normalized_Consumption
14
15 error=as.double(tree_knn[[4]])-testdata$Normalized_Consumption
16 error <- as.data.frame(error)
17 performance <- accuracy(tree_knn[[4]],testdata$Normalized_Consumption)
18 tree_knn[[4]] <- as.data.frame(tree_knn[[4]])
19 colnames(tree_knn[[4]]) <- c("Predicted value")
20 Final_prediction_knn_reg <- cbind(testdata, tree_knn[[4]])
21 Final_prediction_knn_reg$residual <- Final_prediction_knn_reg$Normalized_Consumption -
22   Final_prediction_knn_reg$ Predicted value
23 standard_deviation <- sd(Final_prediction_knn_reg$residual)
24 Final_prediction_knn_reg <- Final_prediction_knn_reg %>% mutate(Outlier_Tag = ifelse(residual >= (2*standard_deviation) , 1, 0))
25
26 write.csv(Final_prediction_knn_reg, "5198_1_knn_reg_Prediction.csv")
27 write.csv(performance, "5198_1_knn_reg_performance.csv")
```

	ME	RMSE	MAE	MPE	MAPE
Test set	0.0008885740011	0.04306537475	0.02618784041	-Inf	Inf

```
> head(Final_prediction_knn_reg)
   TemperatureF hour dayOfWeek Weekday month Holiday Dew_PointF Sea_Level_PressureIn VisibilityMPH
2      33.8     1       2       1     1       1      32.0          29.42           1.9
13     37.4    12       2       1     1       1      35.6          29.30           6.2
17     37.4    16       2       1     1       1      37.4          29.33           2.5
24     37.4    23       2       1     1       1      35.6          29.42           4.3
27     37.4     2       3       1     1       0      35.6          29.44           3.7
33     35.6     8       3       1     1       0      35.6          29.47           1.2
   Normalized_Consumption WindDirDegrees Predicted value      residual Outlier_Tag
2            0.2272727273          150  0.2303030303 -0.003030303030          0
13           0.2363636364          190  0.2272727273  0.009090909091          0
17           0.2181818182          190  0.2303030303 -0.012121212121          0
24           0.2272727273          200  0.2212121212  0.006060606061          0
27           0.2272727273          190  0.2121212121  0.015151515152          0
33           0.2545454545          210  0.1969696970  0.057575757576          0
> |
```

Best Model:

```
> knnReg_RMSE
[1] 0.04306537475
> NEURALnode_RMSE
[1] 0.0430394466
> randomForest_RMSE
[1] 0.04832920999
> regression_tree_rmse
[1] 0.04352705151
> |
```

```
> min(knnReg_RMSE, NEURALnode_RMSE, randomForest_RMSE, regression_tree_rmse)
[1] 0.0430394466
```

Hence least rmse belongs to Neural node

Thus, Neural Node Model is the best model

Classification

Logistic Regression:

```
1 #Some descriptive statistics of Consumption data
2 data_logi <- split_dataset_final[[1]]
3 #data_logi <- data_logi[,c(18,6,8,9,10,12,16,14,19,20,21,22,24,26,28)]
4
5 summary(data_logi)
6
7 data_logi$Base_Hour_Class <- as.factor(data_logi$Base_Hour_Class)
8 table(data_logi$Base_Hour_Class)
9
10 #Construct a logistic regression model for factor Base_Hour_Class using all other variables
11 fit1<-glm(Base_Hour_Class~month+hour+TemperatureF+Dew_PointF+Humidity+Sea_Level_PressureIn+VisibilityMPH+WindDir+
12 +dayOfWeek+Weekday+Holiday,data=data_logi,family=binomial(link='logit'))
13
14 summary(fit1)
15
16 #Construct a logistic regression model for factor Base_Hour_Class
17 fit2 <- glm(Base_Hour_Class ~ month + Sea_Level_PressureIn + Normalized_Consumption + Weekday,
18 data=data_logi, family=binomial(link='logit'))
19 summary(fit2)
20
21 #Parameters of fit2
22 coef(fit2)
23
24 #Predict the probability of the outcome
25 newdata <- data.frame(month=mean(data_logi$month),
26                         Sea_Level_PressureIn=mean(data_logi$Sea_Level_PressureIn),
27                         Normalized_Consumption=mean(data_logi$Normalized_Consumption),
28                         Weekday=mean(data_logi$Weekday))
29 newdata
30 prob <- predict(fit2, newdata=newdata, type="response")
31 prob
32
```

summary(data_logi)

BuildingID	building	meternumb	type	date
Min. :5198	Length:8211	Min. :1	Dist_Heating: 0	Min. :2013-01-01
1st Qu.:5198	Class :character	1st Qu.:1	elect :8211	1st Qu.:2013-03-27
Median :5198	Mode :character	Median :1		Median :2013-06-21
Mean :5198		Mean :1		Mean :2013-06-20
3rd Qu.:5198		3rd Qu.:1		3rd Qu.:2013-09-14
Max. :5198		Max. :1		Max. :2013-12-08
hour	Consumption	dayOfWeek	Weekday	month
Min. : 0.00000	Min. : 0.00000	Min. :0.00000	Min. :0.0000000	Min. : 1.000000
1st Qu.: 5.00000	1st Qu.:22.00000	1st Qu.:1.00000	1st Qu.:0.0000000	1st Qu.: 3.000000
Median :11.00000	Median :24.00000	Median :3.00000	Median :1.0000000	Median : 6.000000
Mean :11.49665	Mean :24.69358	Mean :3.00475	Mean :0.7131896	Mean : 6.158568
3rd Qu.:17.00000	3rd Qu.:26.00000	3rd Qu.:5.00000	3rd Qu.:1.0000000	3rd Qu.: 9.000000
Max. :23.00000	Max. :66.00000	Max. :6.00000	Max. :1.0000000	Max. :12.000000
Base_hour_Flag	Holiday	X..address	area_floor._m.sqr	
Min. :0.0000000	Min. :0.0000000	Metsahovintie 113 :8211	Min. :110	
1st Qu.:0.0000000	1st Qu.:0.0000000	Arkadiankatu 24 : 0	1st Qu.:110	
Median :0.0000000	Median :0.0000000	Otaniementie 17 : 0	Median :110	
Mean :0.2919255	Mean :0.04676653	Betonimiehenkuja 1: 0	Mean :110	
3rd Qu.:1.0000000	3rd Qu.:0.0000000	Betonimiehenkuja 3: 0	3rd Qu.:110	
Max. :1.0000000	Max. :1.0000000	Betonimiehenkuja 5: 0	Max. :110	
	(Other)		: 0	
AirportCode	Normalized_Consumption	Time	TemperatureF	
Length:8211	Min. :0.0000000	Min. :00:00:00	Min. :-22.00000	
Class :character	1st Qu.:0.2000000	1st Qu.:05:50:00	1st Qu.: 30.20000	
Mode :character	Median :0.2181818	Median :11:50:00	Median : 44.60000	
	Mean :0.2244871	Mean :11:50:37	Mean : 44.13898	
	3rd Qu.:0.2363636	3rd Qu.:17:50:00	3rd Qu.: 59.00000	
	Max. :0.6000000	Max. :23:50:00	Max. : 86.00000	
Dew_PointF	Humidity	Sea_Level_PressureIn	VisibilityMPH	
Min. :-27.4000	Min. : 12.00000	Min. :28.74000	Min. : 0.100000	
1st Qu.: 24.8000	1st Qu.: 64.00000	1st Qu.:29.74000	1st Qu.: 6.200000	
Median : 37.4000	Median : 82.00000	Median :29.92000	Median : 6.200000	
Mean : 36.3288	Mean : 76.96395	Mean :29.89768	Mean : 5.880088	
3rd Qu.: 50.0000	3rd Qu.: 93.00000	3rd Qu.:30.09000	3rd Qu.: 6.200000	
Max. : 68.0000	Max. :100.00000	Max. :30.54000	Max. :31.000000	
Wind_Direction	Wind_SpeedMPH	Conditions	WindDirDegrees	base_hr_usage

> table(data_logi\$Base_Hour_Class)

High	Low
4396	3815

#Construct a logistic regression model for factor Base_Hour_Class using all other variables

```

> summary(fit1)

Call:
glm(formula = Base_Hour_Class ~ month + hour + TemperatureF +
    Dew_PointF + Humidity + Sea_Level_PressureIn + VisibilityMPH +
    WindDirDegrees + Wind_SpeedMPH + Normalized_Consumption +
    dayOfWeek + Weekday + Holiday, family = binomial(link = "logit"),
    data = data_logi)

Deviance Residuals:
    Min          1Q          Median         3Q          Max 
-1.8998817 -0.9992677 -0.3381317  0.9487127  3.6045067 

Coefficients:
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) 22.6651607901 3.4284203182 6.61096 0.00000000003818272 ***  
month        0.0377364198 0.0100193800 3.76634 0.00016566 ***  
hour         0.0018339065 0.0037541854 0.48850 0.62519820  
TemperatureF 0.0121815535 0.0185665438 0.65610 0.51175826  
Dew_PointF   -0.0101085397 0.0192513787 -0.52508 0.59952665  
Humidity     -0.0018830176 0.0078822536 -0.23889 0.81118832  
Sea_Level_PressureIn -0.5518168023 0.1069706266 -5.15858 0.00000024882646802 ***  
VisibilityMPH -0.0074748007 0.0107418127 -0.69586 0.48651634  
WindDirDegrees 0.0002039154 0.0002405438 0.84773 0.39659009  
Wind_SpeedMPH -0.0004078618 0.0064862830 -0.06288 0.94986153  
Normalized_Consumption -28.6528311999 1.0535959556 -27.19527 < 0.0000000000000000222 ***  
dayOfWeek      0.0051566266 0.0124140516 0.41539 0.67785916  
Weekday        -0.4215383205 0.0571347357 -7.37797 0.00000000000016072 ***  
Holiday        -0.1015298885 0.1142595591 -0.88859 0.37422353  
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 11341.7178 on 8210 degrees of freedom
Residual deviance: 9440.3229 on 8197 degrees of freedom
AIC: 9468.3229

Number of Fisher Scoring iterations: 5

```

Here, month + Sea_Level_PressureIn + Normalized_Consumption + Weekday are the only important variables with suitable p-values

Hence, Construct a logistic regression model for factor Base_Hour_Class

Summary of selected variables model:

```

Call:
glm(formula = Base_Hour_Class ~ month + Sea_Level_PressureIn +
    Normalized_Consumption + Weekday, family = binomial(link = "logit"),
    data = data_logi)

Deviance Residuals:
    Min          1Q          Median         3Q          Max
-1.8974590 -0.9884293 -0.3397656  0.9591171  3.6630697

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) 19.720544822 2.845829666 6.92963 0.000000000042194 ***
month        0.042281736 0.008377415 5.04711 0.0000004485423604 ***
Sea_Level_PressureIn -0.447278048 0.093443831 -4.78660 0.0000016963217799 ***
Normalized_Consumption -29.454532956 1.022155052 -28.81611 < 0.000000000000000222 ***
Weekday      -0.399792181  0.056487605 -7.07752 0.0000000000014676 ***
---
Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 11341.7178 on 8210 degrees of freedom
Residual deviance: 9468.9183 on 8206 degrees of freedom
AIC: 9478.9183

Number of Fisher Scoring iterations: 5

```

Coefficients:

```

> coef(fit2)
(Intercept)           month   Sea_Level_PressureIn Normalized_Consumption
19.72054482184       0.04228173613      -0.44727804777      -29.45453295563
Weekday                -0.3997921816

```

#Predict the probability of the outcome

```

> newdata <- data.frame(month=mean(data_logi$month),
+                         Sea_Level_PressureIn=mean(data_logi$Sea_Level_PressureIn),
+                         Normalized_Consumption=mean(data_logi$Normalized_Consumption),
+                         Weekday=mean(data_logi$Weekday))
> newdata
  month Sea_Level_PressureIn Normalized_Consumption Weekday
1 6.158567775      29.89768481      0.2244871071  0.7131896237

```

Predict the probability: **0.4282661186**

```

33 #Split the data into training and testing
34 smp_size <- floor(0.80 * nrow(data_logi))
35 set.seed(123)
36 train_ind <- sample(seq_len(nrow(data_logi)), size = smp_size)
37 train <- data_logi[train_ind,]
38 test <- data_logi[-train_ind,]
39 #Build a logistic regression model on the training data
40 fit <- glm(Base_Hour_Class ~ month + Sea_Level_PressureIn + Normalized_Consumption + Weekday,
41             data=train, family=binomial(link="logit"))
42 summary(fit)
43
44 #Run the model on the test set with cutoff value = 0.5
45 test.probs <- predict(fit,test,type="response")
46 pred <- rep("High",length(test.probs))
47 pred[test.probs>=0.5] <- "Low"
48
49 #classification matrix using confusionMatrix() function in caret package
50 library(caret)
51 confusionMatrix(test$Base_Hour_Class, pred)
52 #Generate ROC curve using ROCR package
53 library(ROCR)
54 prediction <- prediction(test.probs, test$Base_Hour_Class)
55 performance <- performance(prediction, measure = "tpr",x.measure = "fpr")
56 plot(performance, main = "ROC curve", xlab = "1-Specificity", ylab = "Sensitivity")
57
58 #Generate cumulative lift curve using lift() function in caret package
59 test$probs = test.probs
60 test$probs = sort(test$probs, decreasing = T)
61 lift <- lift(Base_Hour_Class ~ prob, data = test)
62 lift
63 xyplot(lift, plot = "gain")
64
65

```

#Split the data into training and testing and #Build a logistic regression model on the training data

Summary:

```

Call:
glm(formula = Base_Hour_Class ~ month + Sea_Level_PressureIn +
    Normalized_Consumption + Weekday, family = binomial(link = "logit"),
    data = train)

Deviance Residuals:
    Min      1Q      Median      3Q      Max 
-1.8863156 -0.9865846 -0.3474295  0.9571399  3.6059481 

Coefficients:
              Estimate Std. Error z value Pr(>|z|)    
(Intercept) 21.474539708 3.167970517 6.77864 0.00000000012131 ***
month        0.041005671 0.009378262 4.37242 0.000012287853249 ***
Sea_Level_PressureIn -0.505529769 0.104006077 -4.86058 0.000001170428313 ***
Normalized_Consumption -29.550995928 1.144693674 -25.81564 < 0.00000000000000222 ***
Weekday      -0.389219253 0.063248039 -6.15385 0.000000000756219 ***
---
Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 9067.3857 on 6567 degrees of freedom
Residual deviance: 7574.6660 on 6563 degrees of freedom
AIC: 7584.666

Number of Fisher Scoring iterations: 5

```

#Run the model on the test set with cutoff value = 0.5
#classification matrix using confusionMatrix() function in caret package

Confusion Matrix and Statistics

```
Reference
Prediction High Low
  High  625 238
  Low   233 547

Accuracy : 0.7133293
95% CI  : (0.6907915, 0.7351024)
No Information Rate : 0.5222155
P-Value [Acc > NIR] : < 0.00000000000000022

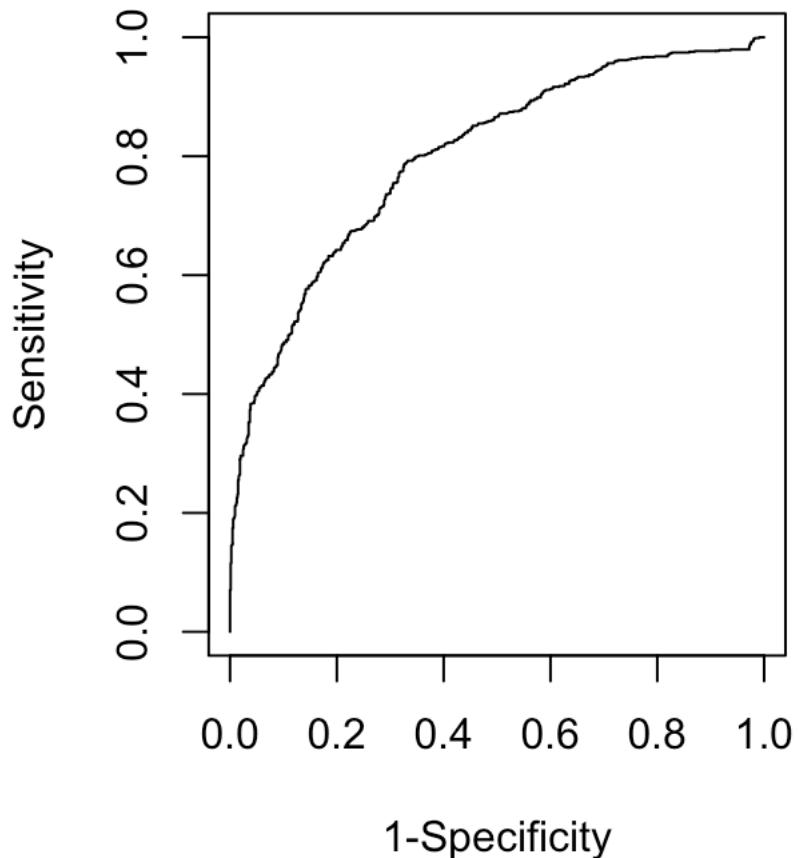
Kappa : 0.4253688
McNemar's Test P-Value : 0.8537701

Sensitivity : 0.7284382
Specificity : 0.6968153
Pos Pred Value : 0.7242178
Neg Pred Value : 0.7012821
Prevalence : 0.5222155
Detection Rate : 0.3804017
Detection Prevalence : 0.5252587
Balanced Accuracy : 0.7126268

'Positive' Class : High
```

```
#Generate ROC curve using ROCR package
```

ROC curve



```
#Generate cumulative lift curve using lift() function in caret package
```

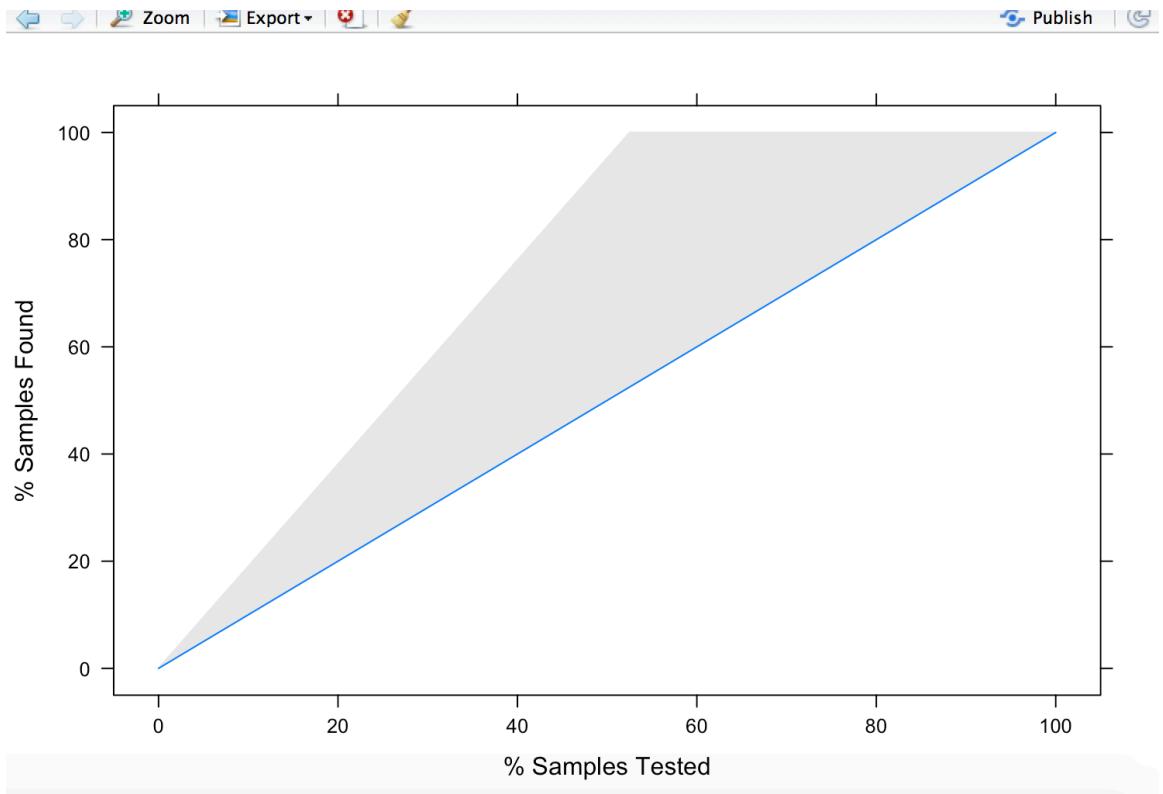
```
> lift
```

Call:

```
lift.formula(x = Base_Hour_Class ~ prob, data = test)
```

Models: prob

Event: High (52.5%)



month	Sea_Level_PressureIn	Normalized_Consumption	Weekday	Base_Hour_Class	pred	OutlierDay	
989	7	29.86	0.2181818182	1	Low	High	False
990	7	29.89	0.2181818182	0	High	Low	False
991	7	29.86	0.1818181818	0	Low	Low	False
992	7	29.89	0.1727272727	0	Low	Low	False
993	7	29.89	0.1909090909	0	Low	Low	False
994	7	29.89	0.1727272727	0	Low	Low	False
995	7	29.92	0.2090909091	0	High	Low	False
996	7	29.92	0.2090909091	0	High	Low	False
997	7	29.95	0.2090909091	0	High	Low	False
998	7	30.01	0.1909090909	0	Low	Low	False
999	7	30.01	0.1818181818	0	Low	Low	False
1000	7	30.06	0.2000000000	1	Low	Low	False
1001	7	30.06	0.1909090909	1	Low	Low	False
1002	7	29.98	0.2000000000	1	Low	Low	False

Random Forest

```
1 set.seed(123)
2 data <- split_dataset_final[[1]]
3 data <- data[,c(18,6,8,9,10,12,7,14,19,20,21,22,24,26,28)]
4 data$Base_Hour_Class <- as.factor(data$Base_Hour_Class)
5 library(randomForest)
6 table(data$Base_Hour_Class)/nrow(data)
7 #High      Low
8 #0.5353793691 0.4646206309
9
10 #53% of the observations has target variable "High" and remaining 46% observations take value "Low".
11
12 #Now, we will split the data sample into development and validation samples.
13
14 sample.ind <- sample(2,
15                       nrow(data),
16                       replace = T,
17                       prob = c(0.6,0.4))
18 cross.sell.dev <- data[sample.ind==1,]
19 cross.sell.val <- data[sample.ind==2,]
20
21 table(cross.sell.dev$Base_Hour_Class)/nrow(cross.sell.dev)
22 #      High      Low
23 #0.5331720105 0.4668279895
24
25 table(cross.sell.val$Base_Hour_Class)/nrow(cross.sell.val)
26 #High      Low
27 #0.5387453875 0.4612546125
28
29 class(cross.sell.dev$Base_Hour_Class)
30 ## [1] "factor"
31
32 varNames <- names(cross.sell.dev)
33 # Exclude ID or Response variable
34 varNames <- varNames[!varNames %in% c("Base_Hour_Class")]
35
36
```

```
table(data$Base_Hour_Class)/nrow(data)
#High      Low
#0.5353793691 0.4646206309
```

#53% of the observations has target variable "High" and remaining 46% observations take value "Low".

#Now, we will split the data sample into development and validation samples.

After split:

```
> table(cross.sell.dev$Base_Hour_Class)/nrow(cross.sell.dev)
```

High	Low
0.5331720105	0.4668279895

```
> table(cross.sell.val$Base_Hour_Class)/nrow(cross.sell.val)
```

High	Low
0.5387453875	0.4612546125

Both development and validation samples have similar target variable distribution.
This is just a sample validation.

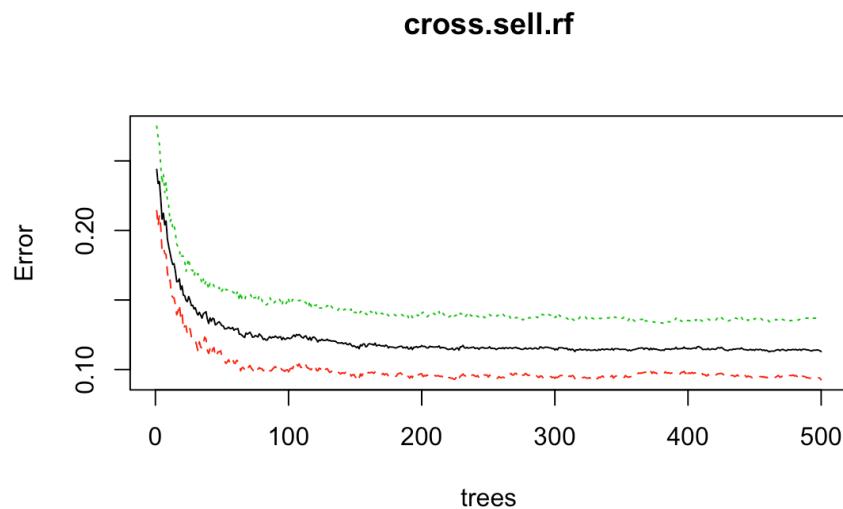
If target variable is factor, classification decision tree is built. We can check the type of response variable.

```
38 # add + sign between exploratory variables
39 varNames1 <- paste(varNames, collapse = "+")
40
41 # Add response variable and convert to a formula object
42 rf.form <- as.formula(paste("Base_Hour_Class", varNames1, sep = " ~ "))
43
44 cross.sell.rf <- randomForest(rf.form,
45                                 cross.sell.dev,
46                                 ntree=500,
47                                 importance=T)
48
49 plot(cross.sell.rf)
50
51 # Variable Importance Plot
52 varImpPlot(cross.sell.rf,
53             sort = T,
54             main="Variable Importance",
55             n.var=5)
56
57 # Variable Importance Table
58 var.imp <- data.frame(importance(cross.sell.rf,
59                                     type=2))
60 # make row names as columns
61 var.imp$Variables <- row.names(var.imp)
62 var.imp[order(var.imp$MeanDecreaseGini,decreasing = T),]
63
64 # Predicting response variable
65 cross.sell.dev$predicted.response <- predict(cross.sell.rf ,cross.sell.dev)
```

```
# Add response variable and convert to a formula object
```

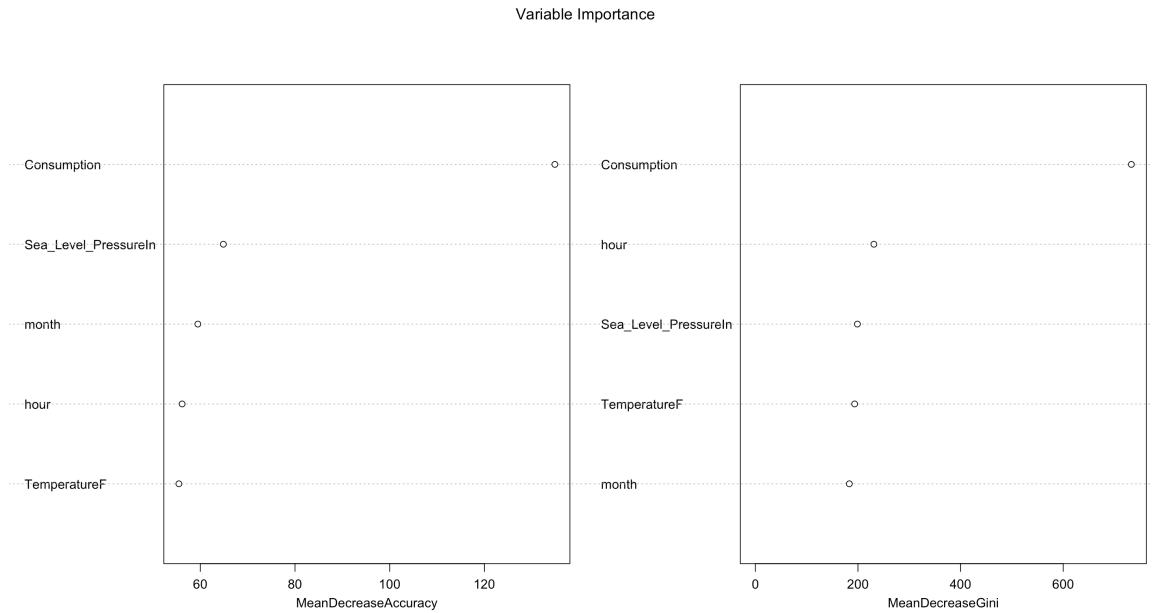
Building Random Forest using R

Now, we have a sample data and formula for building Random Forest model. Let's build 500 decision trees using Random Forest.



500 decision trees or a forest has been built using the Random Forest algorithm based learning. We can plot the error rate across decision trees. The plot seems to indicate that after 100 decision trees, there is not a significant reduction in error rate.

Variable importance plot is also a useful tool and can be plotted using **varImpPlot** function. Top 5 variables are selected and plotted



Predicting response variable

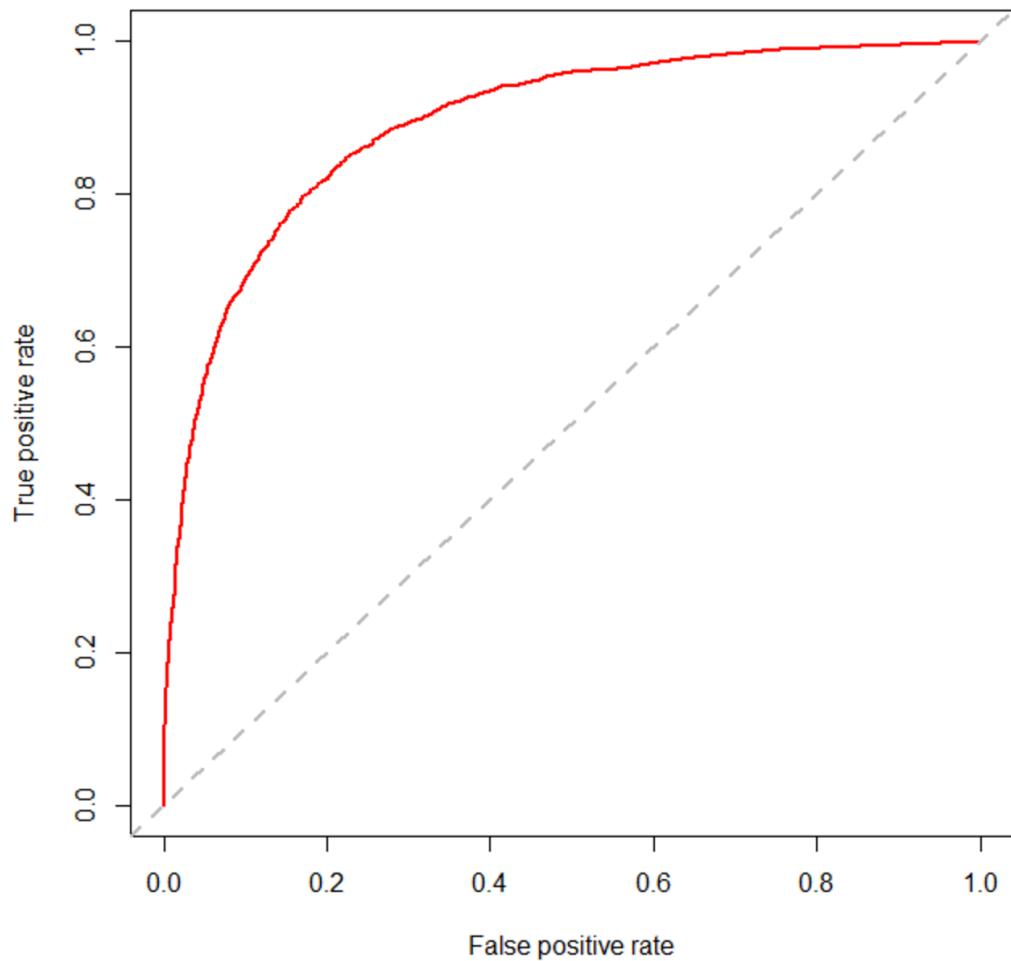
Generic predict function can be used for predicting response variable using Random Forest object.

```
cross.sell.dev$predicted.response <- predict(cross.sell.rf ,cross.sell.dev)
```

Confusion Matrix

confusionMatrix function from caret package can be used for creating confusion matrix based on actual response variable and predicted value.

ROC Curve for Random Forest



```
Reference
Prediction High Low
    High 1565 180
    Low   187 1320

Accuracy : 0.8871464
95% CI : (0.8757678, 0.8978233)
No Information Rate : 0.5387454
P-Value [Acc > NIR] : < 0.0000000000000022

Kappa : 0.7730054
McNemar's Test P-Value : 0.7541307

Sensitivity : 0.8800000
Specificity : 0.8932648
Pos Pred Value : 0.8759124
Neg Pred Value : 0.8968481
Prevalence : 0.4612546
Detection Rate : 0.4059041
Detection Prevalence : 0.4634071
Balanced Accuracy : 0.8866324

'Positive' Class : Low
```

It has 88% accuracy

```

62 var.imp[order(var.imp$MeanDecreaseGini,decreasing = T),]
63 
64 # Predicting response variable
65 cross.sell.dev$predicted.response <- predict(cross.sell.rf ,cross.sell.dev)
66 
67 #Confusion Matrix
68 
69 #confusionMatrix function from caret package can be used for creating confusion matrix
70 
71 # Load Library or packages
72 library(e1071)
73 library(caret)
74 ## Loading required package: lattice
75 ## Loading required package: ggplot2
76 # Create Confusion Matrix
77 confusionMatrix(data=cross.sell.dev$predicted.response,
78                   reference=cross.sell.dev$Base_Hour_Class,
79                   positive='Low')
80 
81 # Predicting response variable
82 cross.sell.val$predicted.response <- predict(cross.sell.rf ,cross.sell.val)
83 
84 # Create Confusion Matrix
85 confusionMatrix(data=cross.sell.val$predicted.response,
86                   reference=cross.sell.val$Base_Hour_Class,
87                   positive='Low')
88

```

```

> head(cross.sell.val)
   TemperatureF hour dayOfWeek Weekday month Holiday Consumption
1       33.8     1        2      1     1      1        25
2       33.8     3        2      1     1      1        25
3       35.6     4        2      1     1      1        25
4       37.4     7        2      1     1      1        27
5       37.4    10        2      1     1      1        26
6       37.4    12        2      1     1      1        26
   area_floor._m.sqr Dew_PointF Humidity Sea_Level_PressureIn VisibilityMPH
1             110     32.0      93        29.42      1.9
2             110     33.8     100        29.36      4.3
3             110     33.8      93        29.33      5.0
4             110     35.6      93        29.30      6.2
5             110     37.4     100        29.30      5.0
6             110     35.6      93        29.30      6.2
   Wind_SpeedMPH WindDirDegrees Base_Hour_Class
1           17.3          150        Low
2           15.0          150        Low
3           15.0          150        Low
4           12.7          180       High
5           16.1          190       High
6           15.0          190       High
   cross.sell.val$predicted.response OutlierDay
1                               Low      True
2                               Low      True
3                               Low      True
4                               High     True
5                               High     True
6                               High     True

```

KNN Classification

```
1 #look up the structure of the dataset#
2 data_knn <- split_dataset_final$`5198.1` 
3 data_knn <- data_knn[,c(18,6,8,9,10,12,16,14,19,20,21,22,24,26,28)]
4
5 str(data_knn)
6 head(data_knn)
7 table(data_knn$Base_Hour_Class)
8
9 #Mix up dataset#
10 set.seed(9850)
11 group<- runif(nrow(data_knn))
12 data_knn<- data_knn[order(group),]
13 summary(data_knn[, c(1:14)])
14
15 #Normalization#
16 normalize<-function(x){return((x-min(x))/(max(x)-min(x)))}
17 data_knn_n<-as.data.frame(lapply(data_knn[,c(1:7,9:14)],normalize))
18 summary(data_knn_n)
19
20 #Separate training&testing dataset
21 data_knn_n_train<-data_knn_n[1:7000, ]
22 data_knn_n_test<-data_knn_n[7001:8211, ]
23 data_knn_train_target<-data_knn[1:7000,15]
24 data_knn_test_target<-data_knn[7001:8211,15]
25
26 # fit in knn algorithm
27 require(class)
28 m1<- knn(train= data_knn_n_train, test= data_knn_n_test, cl=as.factor(data_knn_train_target[['Base_Hour_Class']]), k=92)
29 table(as.factor(data_knn_test_target[['Base_Hour_Class']]),m1)
30
31 #install.packages("gmodels")
32
33 library(gmodels)
34 #result and comparison
35
36 CrossTable(x=as.factor(data_knn_test_target[['Base_Hour_Class']]), y=m1, prop.chisq = FALSE)
37
```

> **table(data_knn\$Base_Hour_Class)**

High	Low
4396	3815

#Mix up dataset#

Since it is a very nice organized dataset, we need to mix the entire dataset up. To do this, we have to create a 150 random number dataset and make the original dataset re-organized follow the order of the random number dataset.

```
> summary(data_knn[, c(1:14)])
   TemperatureF      hour      dayOfWeek      Weekday      month      Holiday
Min. :-22.00000  Min. : 0.00000  Min. : 0.000000  Min. : 0.000000  Min. : 1.000000  Min. :-0.00000000
1st Qu.: 30.20000  1st Qu.: 5.00000  1st Qu.:1.00000  1st Qu.:0.000000  1st Qu.: 3.000000  1st Qu.:0.00000000
Median : 44.60000  Median :11.00000  Median :3.00000  Median :1.000000  Median : 6.000000  Median :0.00000000
Mean   : 44.13898  Mean   :11.49665  Mean   :3.00475  Mean   :0.7131896 Mean   : 6.158568  Mean   :0.04676653
3rd Qu.: 59.00000  3rd Qu.:17.00000  3rd Qu.:5.00000  3rd Qu.:1.000000  3rd Qu.: 9.000000  3rd Qu.:0.00000000
Max.   : 86.00000  Max.   :23.00000  Max.   :6.00000  Max.   :1.0000000 Max.   :12.000000  Max.   :1.00000000
Normalized_Consumption area_floor_m.sqr Dew_PointF      Humidity      Sea_Level_PressureIn VisibilityMPH
Min. : 0.0000000  Min. :110      Min. :-27.4000  Min. : 12.00000  Min. :28.74000  Min. : 0.100000
1st Qu.:0.2000000  1st Qu.:110    1st Qu.: 24.8000  1st Qu.: 64.00000  1st Qu.:29.74000  1st Qu.: 6.200000
Median :2.1818182  Median :110    Median : 37.4000  Median : 82.00000  Median :29.92000  Median : 6.200000
Mean   :0.2244871  Mean   :110    Mean   : 36.3288  Mean   : 76.96395 Mean   :29.89768  Mean   : 5.880088
3rd Qu.:0.2363636  3rd Qu.:110   3rd Qu.: 50.0000  3rd Qu.: 93.00000  3rd Qu.:30.09000  3rd Qu.: 6.200000
Max.   : 0.6000000  Max.   :110    Max.   : 68.0000  Max.   :100.00000 Max.   :30.54000  Max.   :31.000000
Wind_SpeedMPH      WindDirDegrees
Min. : 1.200000  Min. : 0.000
1st Qu.: 4.600000  1st Qu.: 90.000
Median : 8.100000  Median :200.000
Mean   : 8.250432  Mean   :183.393
3rd Qu.:11.500000  3rd Qu.:270.000
Max.   :33.400000  Max.   :360.000
```

#Normalization#

```
> summary(data_knn_n)
   TemperatureF      hour      dayOfWeek      Weekday      month      Holiday
Min. : 0.0000000  Min. : 0.0000000  Min. : 0.0000000  Min. : 0.0000000  Min. : 0.00000000
1st Qu.:0.4833333  1st Qu.:0.2173913  1st Qu.:0.1666667  1st Qu.:0.0000000  1st Qu.:0.1818182  1st Qu.:0.00000000
Median :0.6166667  Median :0.4782609  Median :0.5000000  Median :1.0000000  Median :0.4545455  Median :0.00000000
Mean   :0.6123980  Mean   :0.4998544  Mean   :0.5007916  Mean   :0.7131896  Mean   :0.4689607  Mean   :0.04676653
3rd Qu.:0.7500000  3rd Qu.:0.7391304  3rd Qu.:0.8333333  3rd Qu.:1.0000000  3rd Qu.:0.7272727  3rd Qu.:0.00000000
Max.   : 1.0000000  Max.   :1.0000000  Max.   :1.0000000  Max.   :1.0000000  Max.   :1.00000000  Max.   :1.00000000
Normalized_Consumption Dew_PointF      Humidity      Sea_Level_PressureIn VisibilityMPH
Min. : 0.0000000  Min. :0.0000000  Min. : 0.0000000  Min. : 0.0000000  Min. : 0.00000000
1st Qu.:0.3333333  1st Qu.:0.5471698  1st Qu.: 0.5909091  1st Qu.: 0.5555556  1st Qu.:0.1974110
Median :0.3636364  Median :0.6792453  Median :0.7954545  Median :0.6555556  Median :0.1974110
Mean   :0.3741452  Mean   :0.6680168  Mean   : 0.7382267  Mean   : 0.6431582  Mean   : 0.1870579
3rd Qu.:0.3939394  3rd Qu.:0.8113208  3rd Qu.: 0.9204545  3rd Qu.: 0.7500000  3rd Qu.:0.1974110
Max.   : 1.0000000  Max.   :1.0000000  Max.   :1.0000000  Max.   :1.0000000  Max.   :1.00000000
Wind_SpeedMPH      WindDirDegrees
Min. : 0.0000000  Min. : 0.000000
1st Qu.: 0.1055901  1st Qu.: 0.2500000
Median : 0.2142857  Median : 0.5555556
Mean   : 0.2189575  Mean   : 0.5094250
3rd Qu.:0.3198758  3rd Qu.: 0.7500000
Max.   : 1.0000000  Max.   : 1.0000000
```

#Separate training&testing dataset

fit in knn algorithm

We use knn() function to do the training, we need to fit in the training data frame, the test data frame and the training target variables. before that we have to choose a k value, as k is a place holder for how many nearest neighbor that you want the algorithm to use. The rule of thumb is to take the square root of the total number of observations you have, in this case should be 92

To see how well the model predicted, we use table() to see the difference between predict result and what species they actually are.

```
> table(as.factor(data_knn_test_target[['Base_Hour_Class']]),m1)
m1
```

	High	Low
High	463	170
Low	221	357

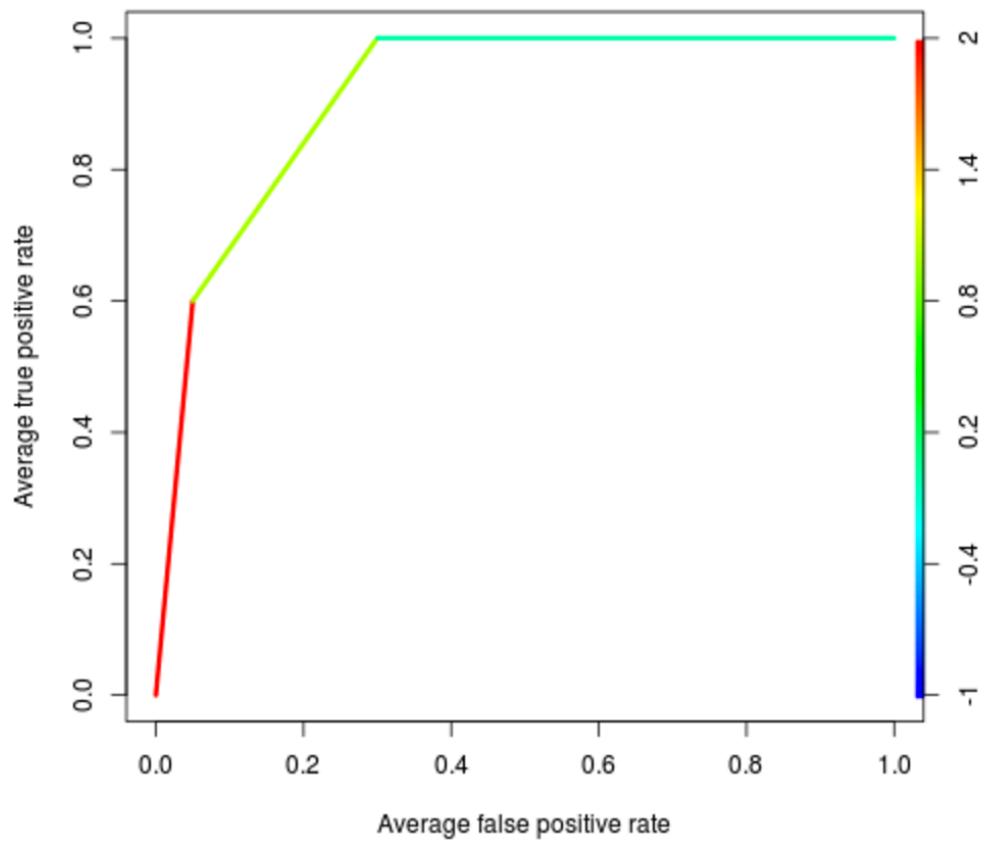
```
> CrossTable(x=as.factor(data_knn_test_target[['Base_Hour_Class']]), y=m1, prop.chisq = FALSE)
```

Cell Contents	
	N
N / Row Total	
N / Col Total	
N / Table Total	

Total Observations in Table: 1211

as.factor(data_knn_test_target[["Base_Hour_Class"]])	m1			Row Total
	High	Low	Column Total	
High	463	170	633	633
	0.731	0.269	0.523	
	0.677	0.323		
	0.382	0.140		
Low	221	357	578	578
	0.382	0.618	0.477	
	0.323	0.677		
	0.182	0.295		
Column Total	684	527	1211	1211
	0.565	0.435		

```
> Final_data[767:770,]
   TemperatureF hour dayOfWeek Weekday month Holiday Normalized_Consumption
767        39.2     1       5      1     4       0           0.1818181818
768        44.6     6       0      0    11       0           0.409090909091
769         6.8     0       1      1     3       0           0.1818181818
770        23.0    23       0      0    12       0           0.0000000000
   area_floor._m.sqr Dew_PointF Humidity Sea_Level_PressureIn VisibilityMPH
767            110     39.2     100        29.53      1.7
768            110     39.2      81        29.39      6.2
769            110     -5.8      56        30.36      6.2
770            110     19.4      86        29.86      6.2
   Wind_SpeedMPH WindDirDegrees Base_Hour_Class Prediction_value OutlierDay
767        10.4          200        Low        High    False
768        15.0          250        High       Low    False
769         9.2          50        Low       Low    True
770         5.8          20        Low       Low    True
```



Best Model: Random Forest with accuracy: 88%

Clustering

```
1  data <- split_dataset_final[[1]]
2  data <- data[,c(18,6,8,9,10,19,20,21,22,24,26,16)]
3  json_data <- fromJSON(paste(readLines("config.json"), collapse=""))
4  #nstart tells how many times algorithm
5  km.out <- kmeans(data,as.numeric(json_data$Number_of_clusters),nstart = as.numeric(json_data$nstart))
6  names(km.out)
7
8  #kmeans results
9  km.out$cluster
10
11 #cluster tagging
12 View(cbind(data,km.out$cluster))
13 View(data)
14 #k-mean results with k=3
15
16 #Scatter plot matrix
17 plot(data, col=(km.out$cluster), main = "K-mean results")
18
19 error.freq.long <- gather(data, tag, freq)
20
21
22 ##### Hierarchical clustering
23 rough_data <- sample(1:nrow(data),round(0.60*nrow(data)))
24
25 k_mean_data <- data[rough_data,]
26 kmeanstestdata <- data[-rough_data,]
27
28 #Scaling the data
29 data=scale(k_mean_data)
30 # Complete linkage type
31 hc.complete=hclust(dist(data),method="complete")
32 # Average linkage type
33 hc.average=hclust(dist(data),method="average")
34
```

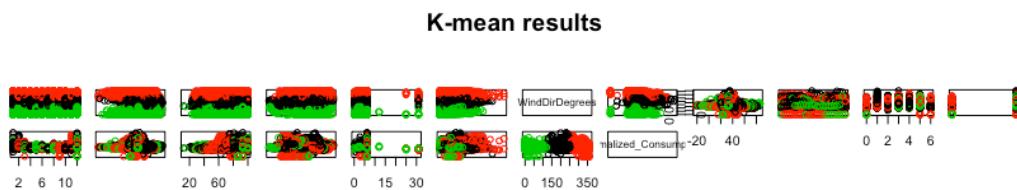
```

26 kmeanstestdata <- data[-rough_data,]
27
28 #Scaling the data
29 data=scale(k_mean_data)
30 # Complete linkage type
31 hc.complete=hclust(dist(data),method="complete")
32 # Average linkage type
33 hc.average=hclust(dist(data),method="average")
34
35 par(mfrow=c(1,2)) #Plotting in a matrix form
36 plot(hc.complete,main='Complete')
37
38
39 ##### Bend Graph
40 rough_data <- sample(1:nrow(data),round(0.75*nrow(data)))
41 k_mean_data <- data[rough_data,]
42 nrow(k_mean_data)
43 names(data)
44 data=scale(k_mean_data)
45 #Scaling the data
46 wss <- (nrow(data)-1)*sum(apply(data,2,var))
47 for(i in 2:15){
48   wss[i] <- sum(kmeans(data,centers = i)$withinss)
49 }
50
51 plot(1:15, wss,type="b",xlab = "Number of clusters", ylab="Within groups sum of squares")
52

```

#Scatter plot matrix

plot(data, col=(km.out\$cluster), main = "K-mean results")



```
> error.freq.long
# A tibble: 98,532 × 2
      tag   freq
      <chr> <dbl>
1 TemperatureF 35.6
2 TemperatureF 33.8
3 TemperatureF 33.8
4 TemperatureF 33.8
5 TemperatureF 35.6
6 TemperatureF 35.6
7 TemperatureF 37.4
8 TemperatureF 37.4
9 TemperatureF 37.4
10 TemperatureF 37.4
# ... with 98,522 more rows
```

Complete linkage type

```
> hc.complete
```

Call:

```
hclust(d = dist(data), method = "complete")
```

Cluster method : complete

Distance : euclidean

Number of objects: 4927

Average linkage type

```
> hc.average
```

Call:

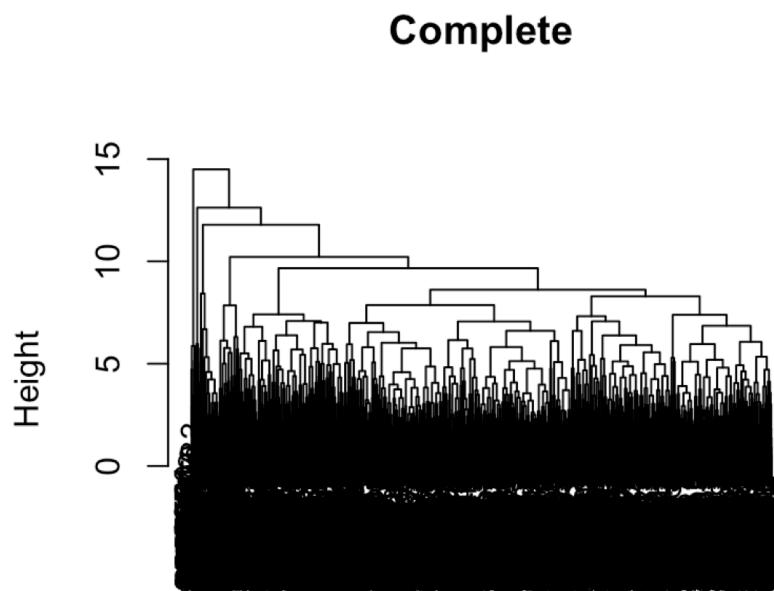
```
hclust(d = dist(data), method = "average")
```

Cluster method : average

Distance : euclidean

Number of objects: 4927

```
#Plotting in a matrix form  
#### Hierarchical clustering
```



```
dist(data)  
hclust (*, "complete")
```

